

Open Razzmatazz Laboratory (OrzLab)

深入淺出 Hello World – Part II

探索記憶體模型與系統呼叫

Nov 25, 2006

Jim Huang(黃敬群 /jserv)
Email: <jserv.tw@gmail.com>
Blog: <http://blog.linux.org.tw/jserv/>

注意

- 本議程針對 x86 硬體架構，至於 ARM 與 MIPS 架構，請另行聯絡以作安排
- 簡報採用創意公用授權條款 (Creative Commons License: **Attribution-ShareAlike**) 發行
- 議程所用之軟體，依據個別授權方式發行
- 系統平台
 - Ubuntu feisty (development branch)
 - Linux kernel 2.6.17-9
 - gcc 4.1.2 (pre-release)
 - glibc 2.5



大綱

- 探索 Linux 記憶體模型
- 深入 syscall
- 再探動態連結

探索 Linux 記憶體模型

- Memory Section/Region
- Linux Memory Allocation
- i386 stack/register/call

Memory section (1)

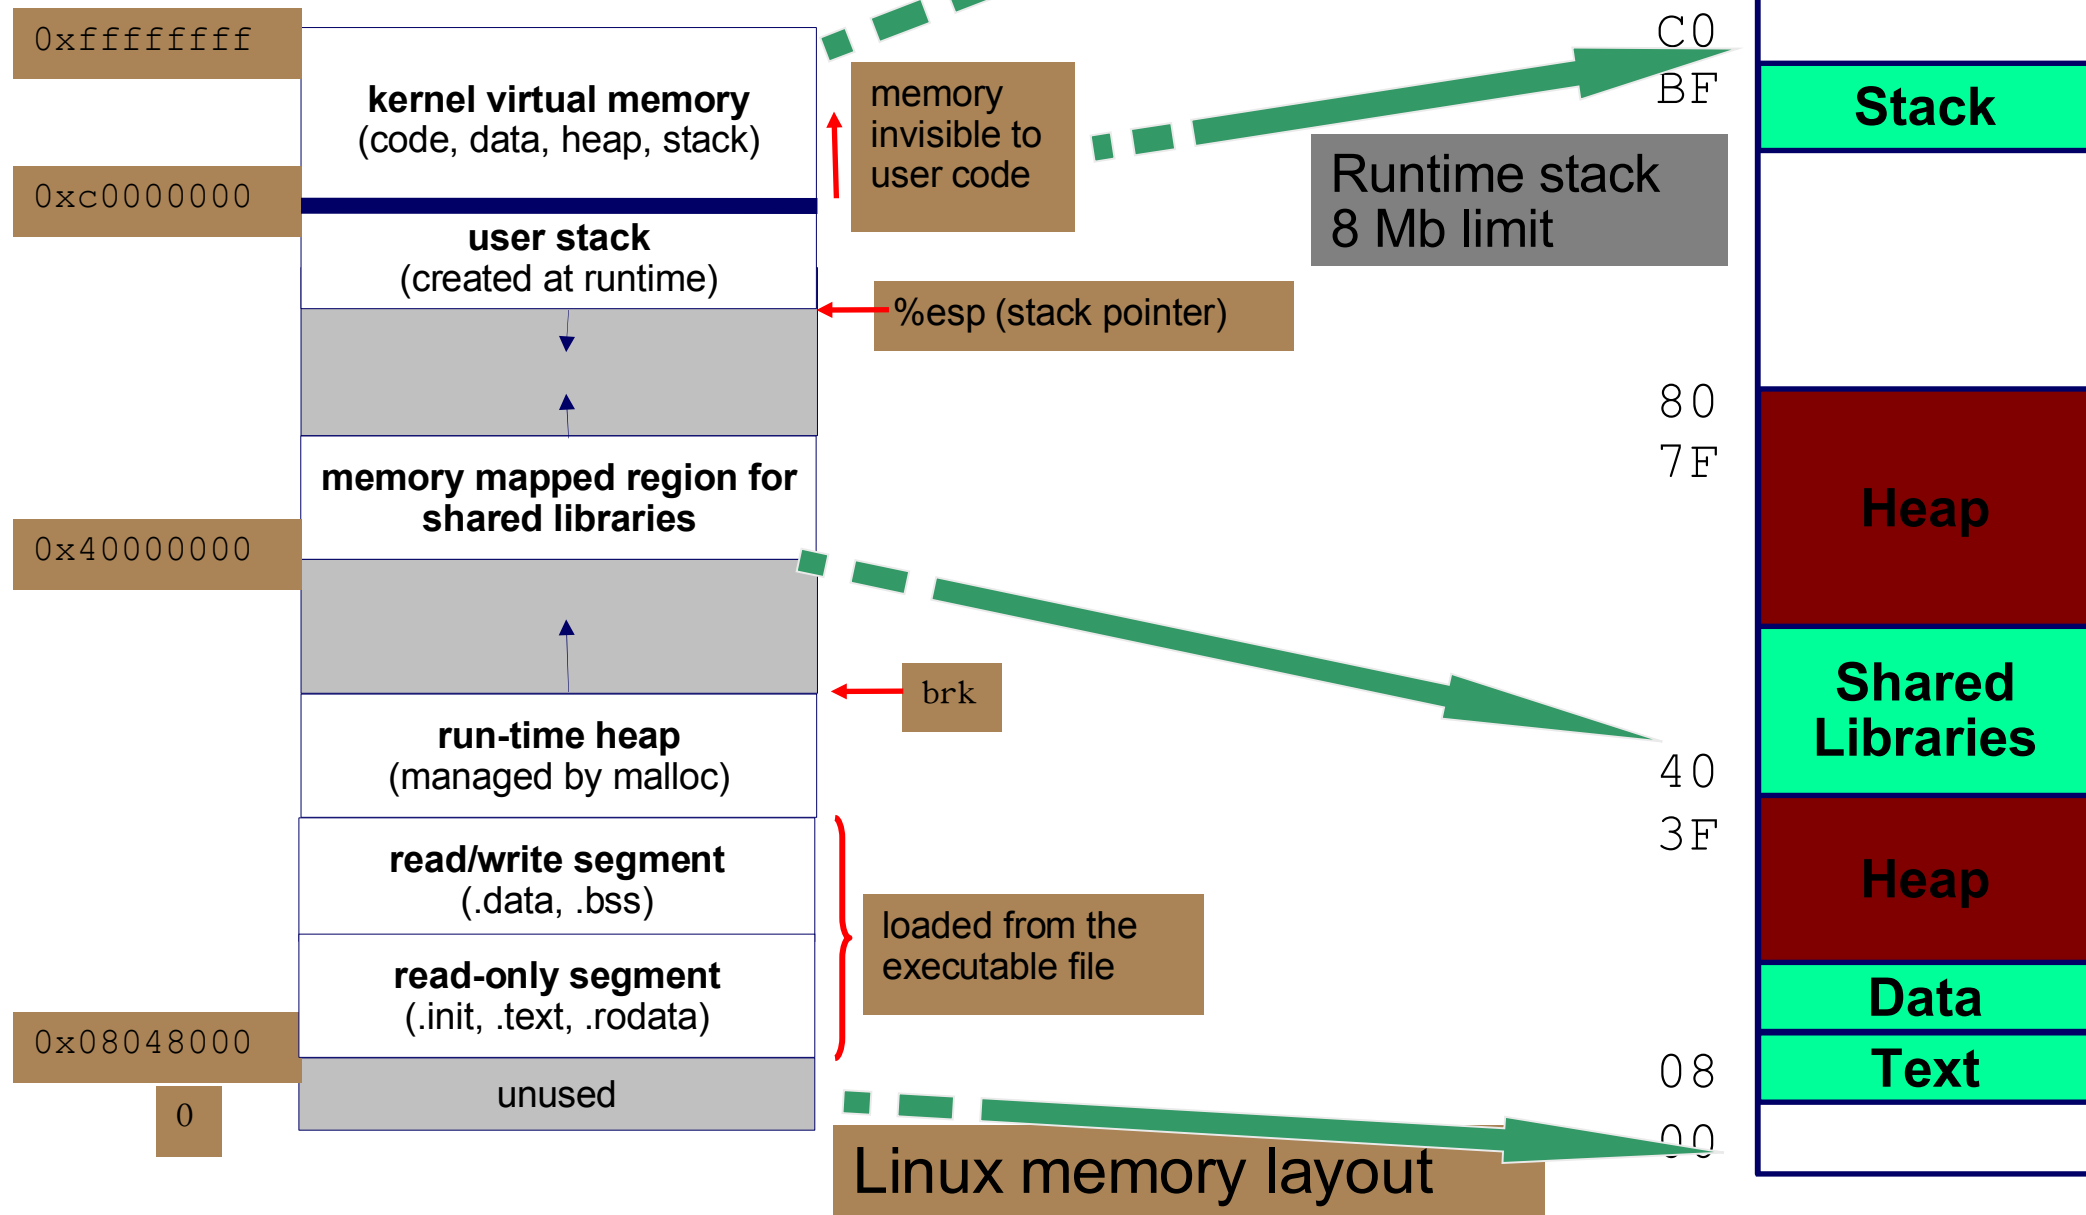
- 概念：Modular programming
- Module 就是特定 I/O 介面的黑盒子，由 re-locatable object code 組成
- Assembler 與 Linker 層面來說，程式由若干 **sections** (absolute or relocatable blocks of code or data) 集合而成
- 來自不同 module、卻有相同名稱的 sections，特稱為“**partial sections**”
- Linker 的職責就是將所有 partial sections 結合為 sections

Memory section (2)

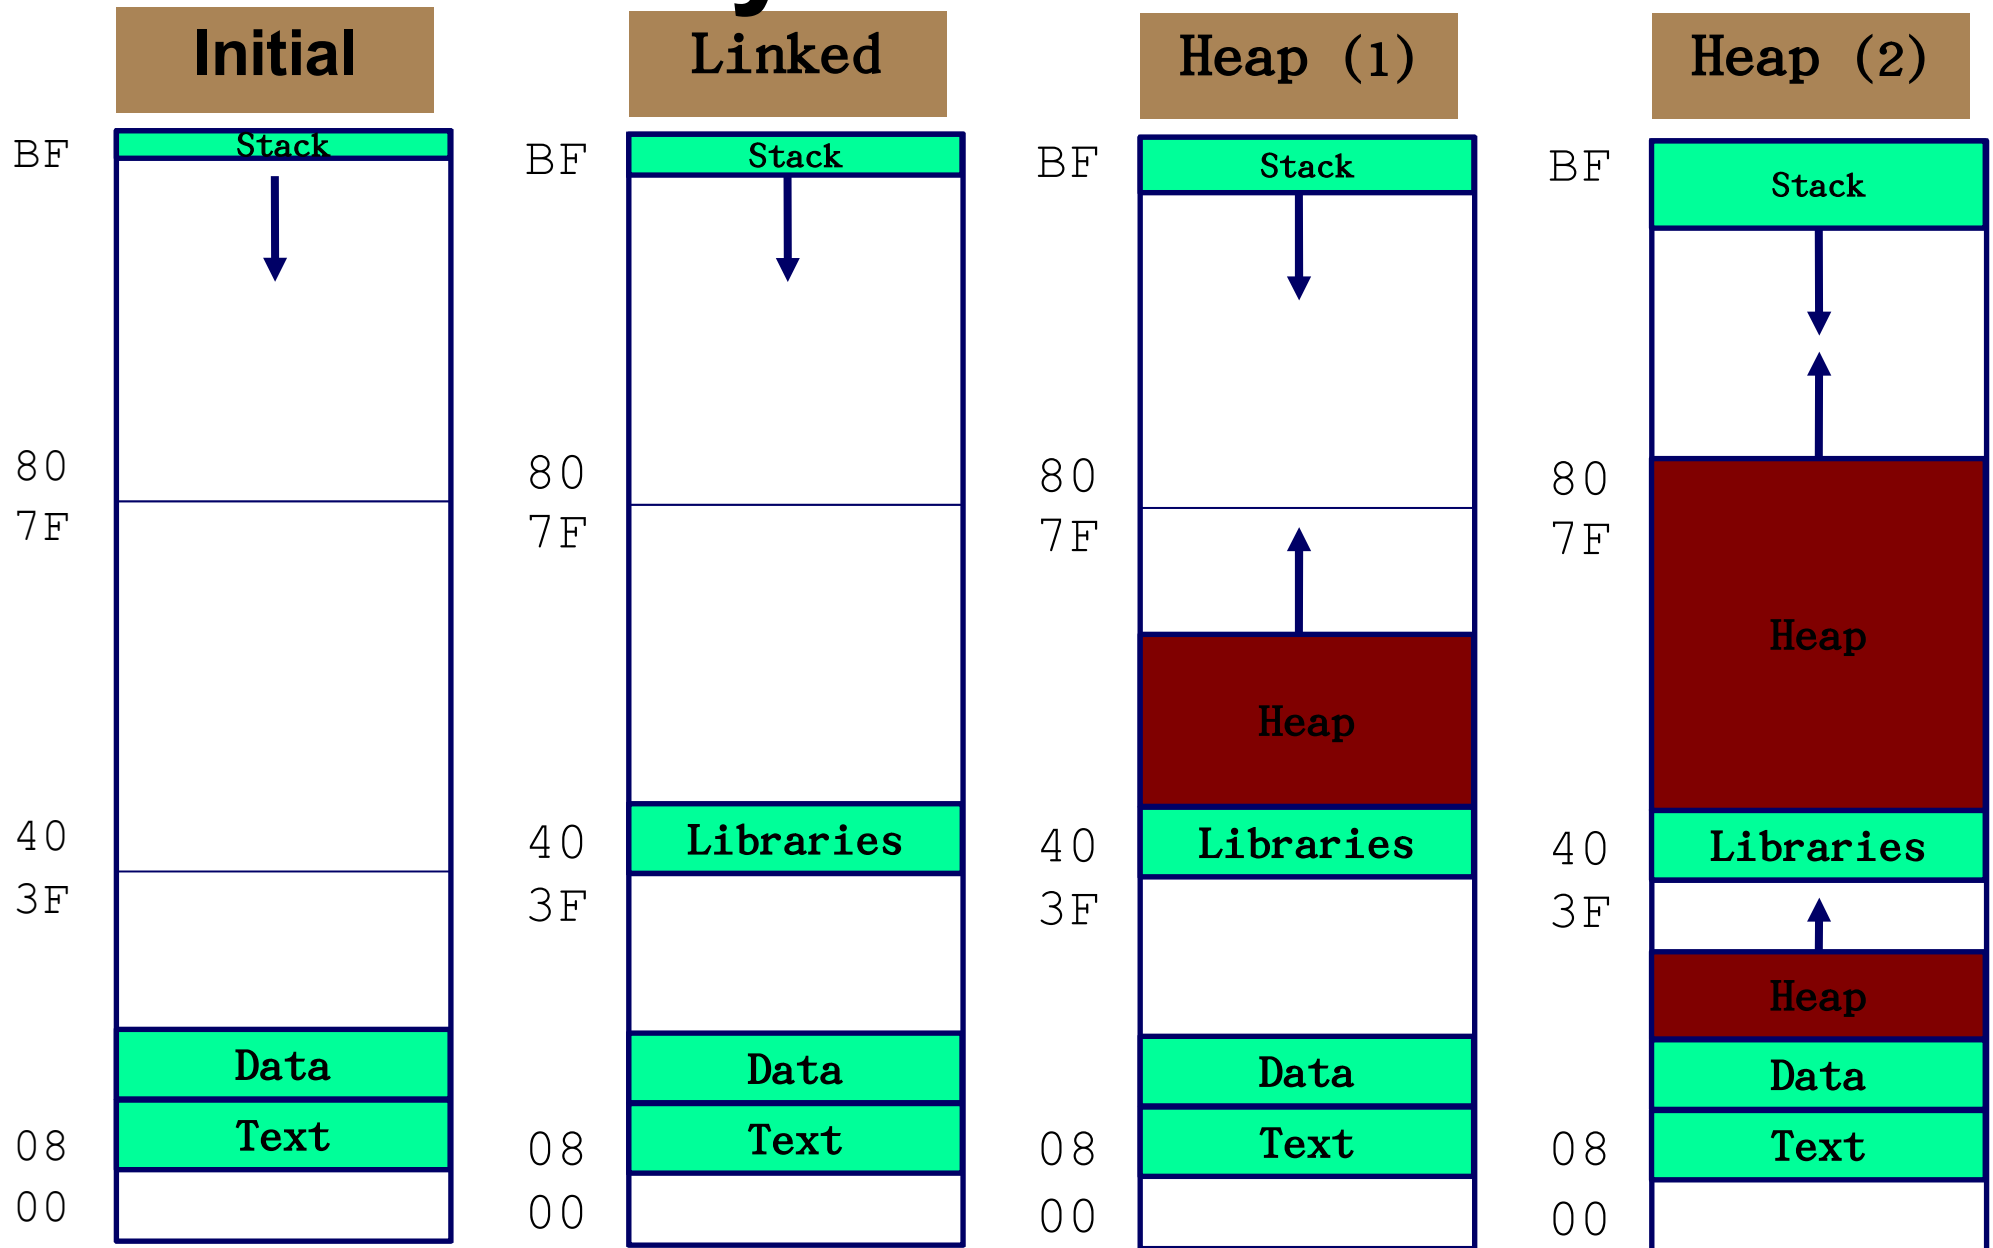
- Absolute sections
 - 需要定位於特定的 memory address
 - 不可結合到其他 section
- Module
 - 由一個或多個 code sections 所組成
- Program
 - 包含單一 absolute module ，並整合其他 absolute/relocatable section
- Linker
 - all external references to symbols in other modules are **resolved** by the linker/locator – the end product is a single absolute object module (*the “program”*)

每個 process 都有獨自的 address space

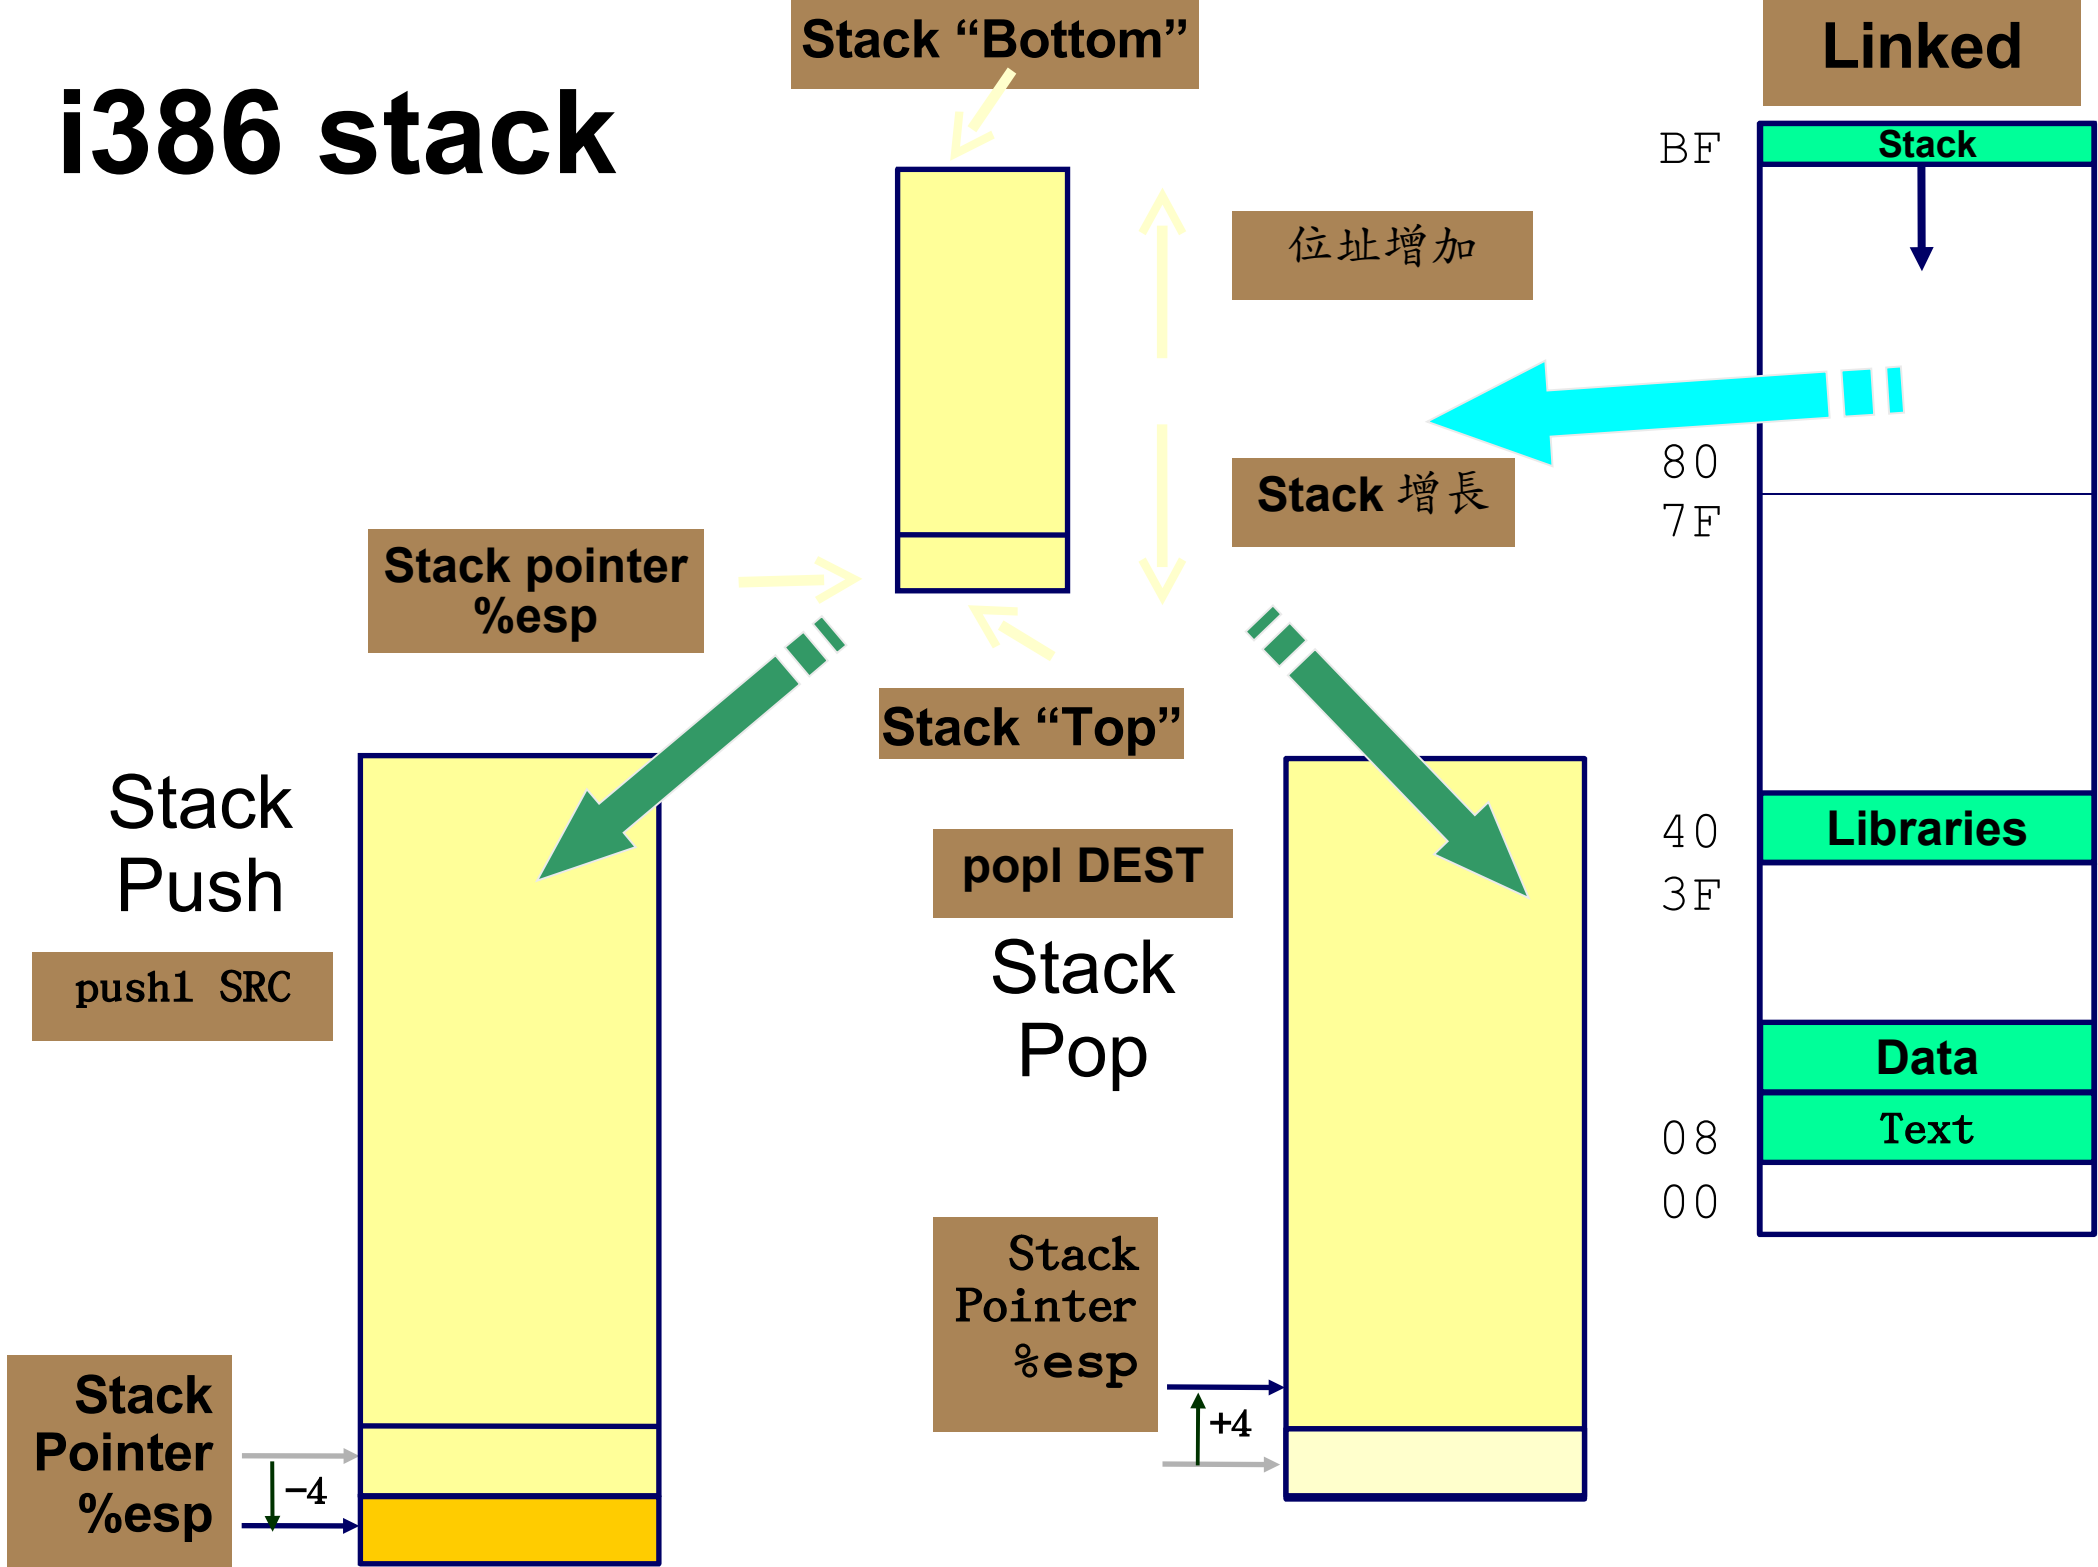
Address Space



Linux Memory Allocation



i386 stack



i386/Linux register

也作為保存 return address 使用

```
#include <stdio.h>
char message[] = "Hello, world!\n";
int main(void)
{
    long _res;
    __asm__ volatile (
        "int $0x80"
        : "=a" (_res)
        : "a" ((long) 4),
          "b" ((long) 1),
          "c" ((long) message),
          "d" ((long) sizeof(message)));
    return 0;
}
```

```
.text
message:
.ascii "Hello World!\0"
.align 4
.globl main
main:
    pushl %ebp
    movl %esp,%ebp
    pushl $message
    call puts
    addl $4,%esp
    xorl %eax,%eax
    movl %ebp,%esp
    popl %ebp
    ret
```

Caller-Save
Temporaries

Callee-Save
Temporaries

Special

%eax

%edx

%ecx

%ebx

%esi

%edi

%esp

%ebp

i386 call(1)

Disassembly of section .text:

_start 為該 Image 的 entry point

080482b0 <_**start**>:

80482b0: 31 ed xor %ebp,%ebp
80482b2: 5e pop %esi
80482b3: 89 e1 mov %esp,%ecx
80482b5: 83 e4 f0 and \$0xffffffff0,%esp
80482b8: 50 push %eax

08048280 <__libc_start_main@plt>:

8048280: ff 25 44 95 04 08 jmp *0x8049544
8048286: 68 00 00 00 00 push \$0x0
804828b: e9 e0 ff ff ff jmp 8048270 <_init+0x18>

80482c5: 51 push %ecx
80482c6: 56 push %esi
80482c7: 68 50 83 04 08 push \$0x8048350
80482cc: e8 af ff ff ff call 8048280 <__libc_start_main@plt>

80482d1: f4 hlt
80482d2: 90 nop
80482d3: 90 nop

◆ call LABEL

- 使用 stack 來實現 procedure call
- 先 PUSH 返回位址，然後 JUMP 到 LABEL

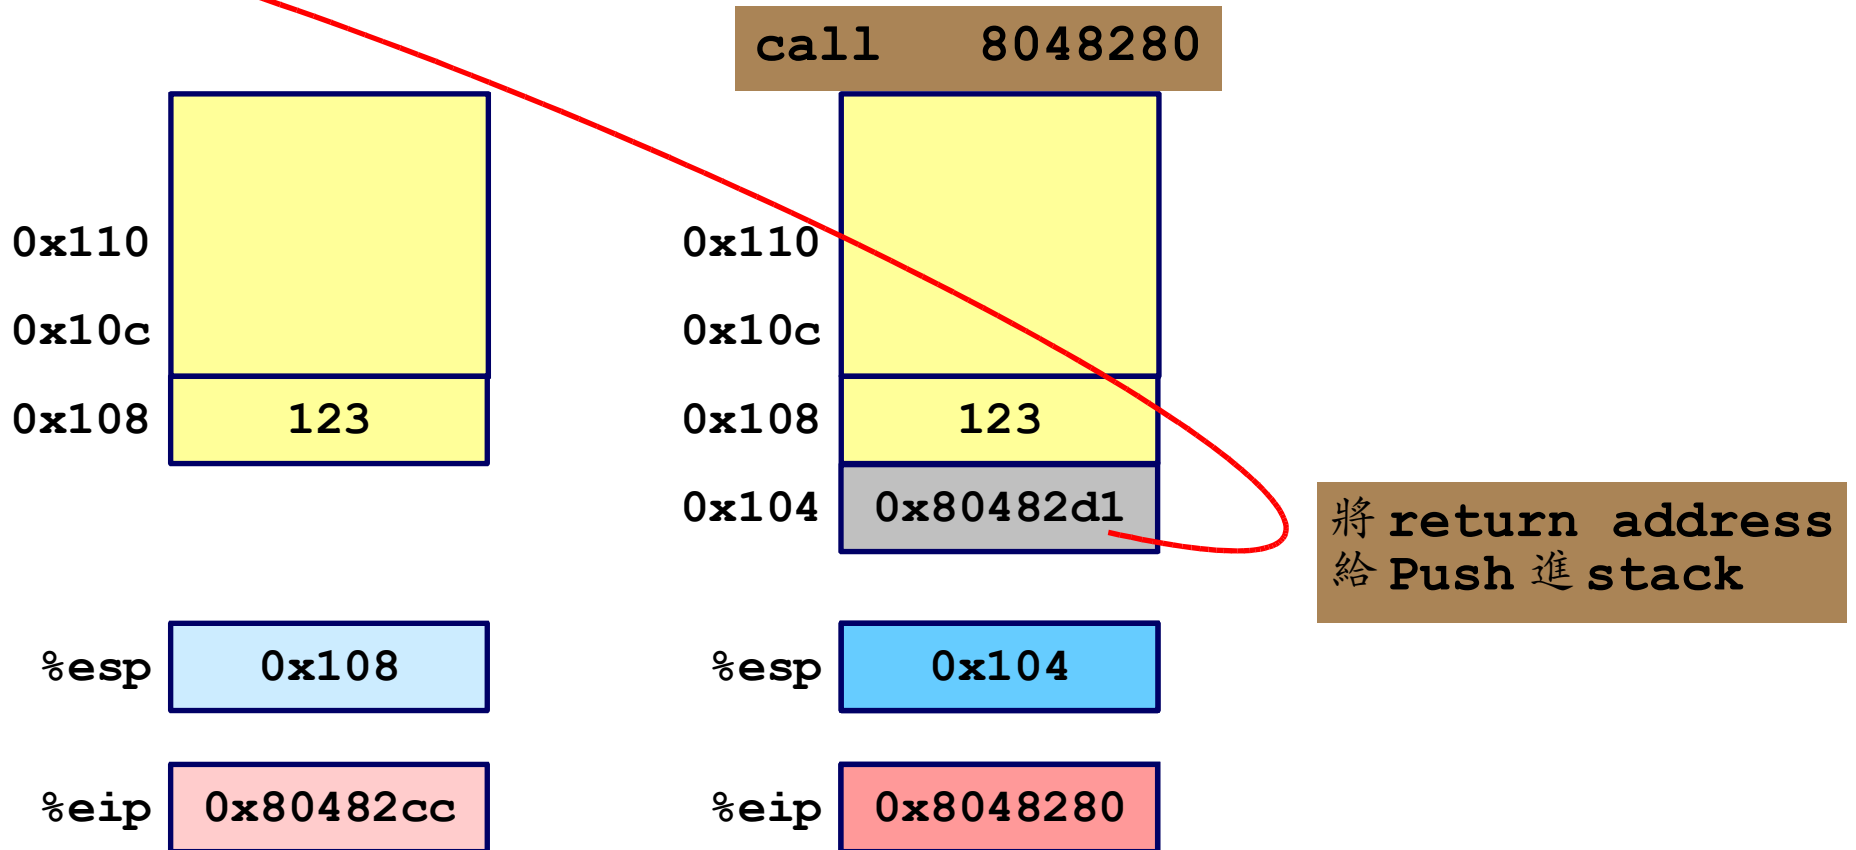
i386 call(2)

Disassembly of section .text:

080482b0 <_start>:

...

```
80482cc:    e8 af ff ff ff    call 8048280 <__libc_start_main@plt>
80482d1:    f4              hlt
```



將 **return address**
給 **Push** 進 **stack**

Programmer Counter

GCC Inner Function

```
$ cat hello.c
#include <stdio.h>

void invoke(void (*func)()) {
    func();
}

void outer() {
    char *hello = "Hello World!\n";
    void inner() {
        printf("%s", hello);
    }
    invoke(inner);
}

int main() {
    outer();
}

$ ./hello
Hello World!
```

call ptr

outer()

invoke()

inner()

將 return address
給 Push 進 stack

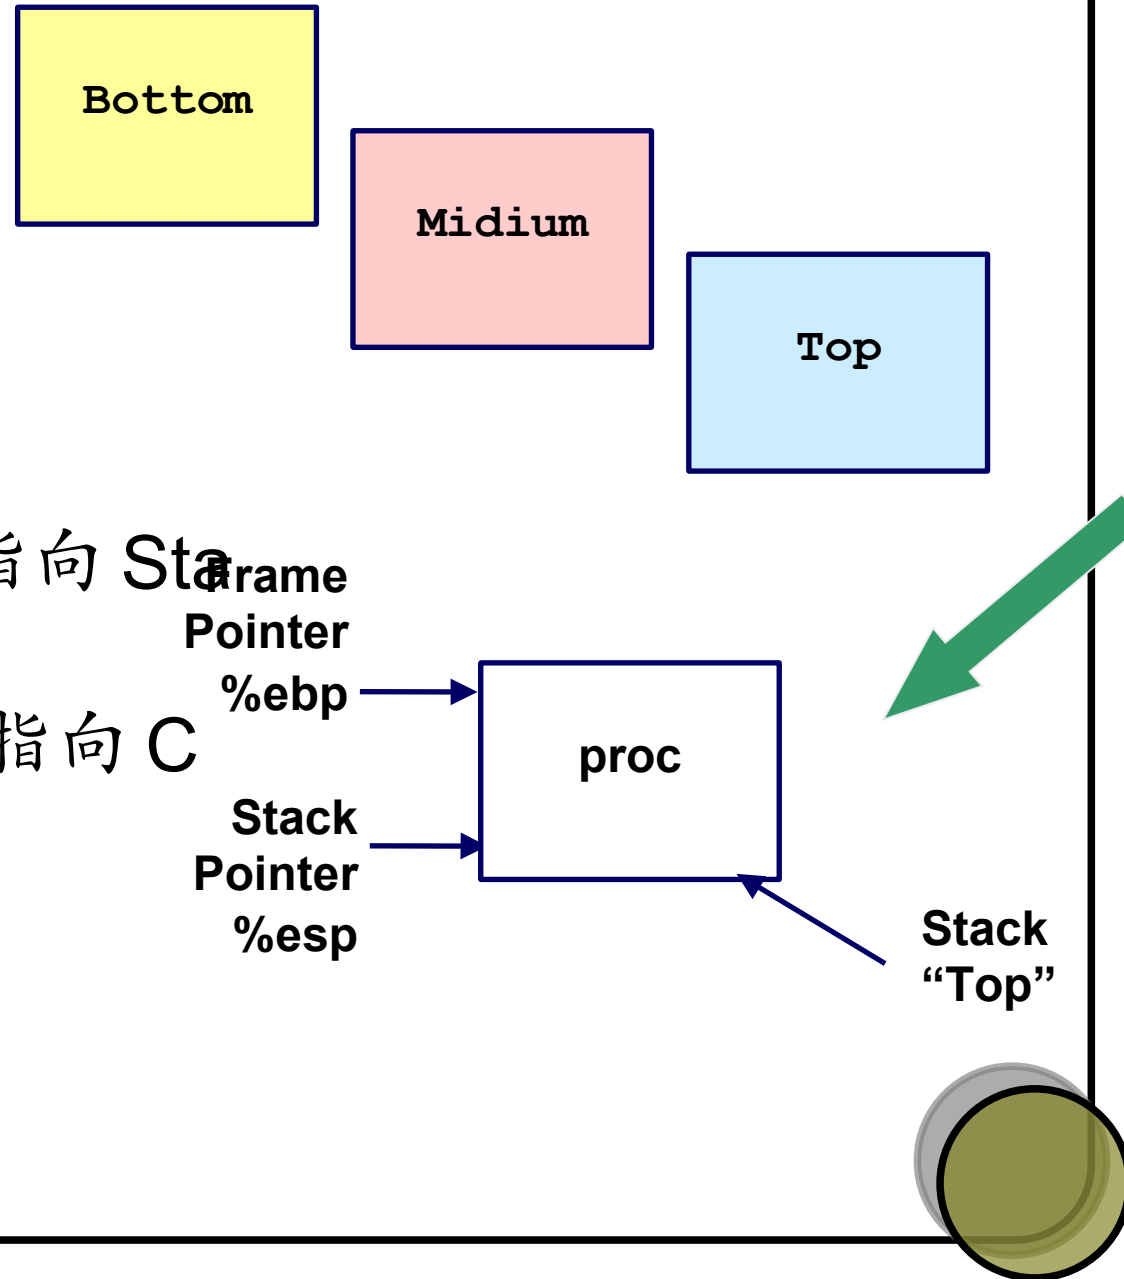
Stack Frame

- 要素

- Local variables
- Return information
- Temporary space

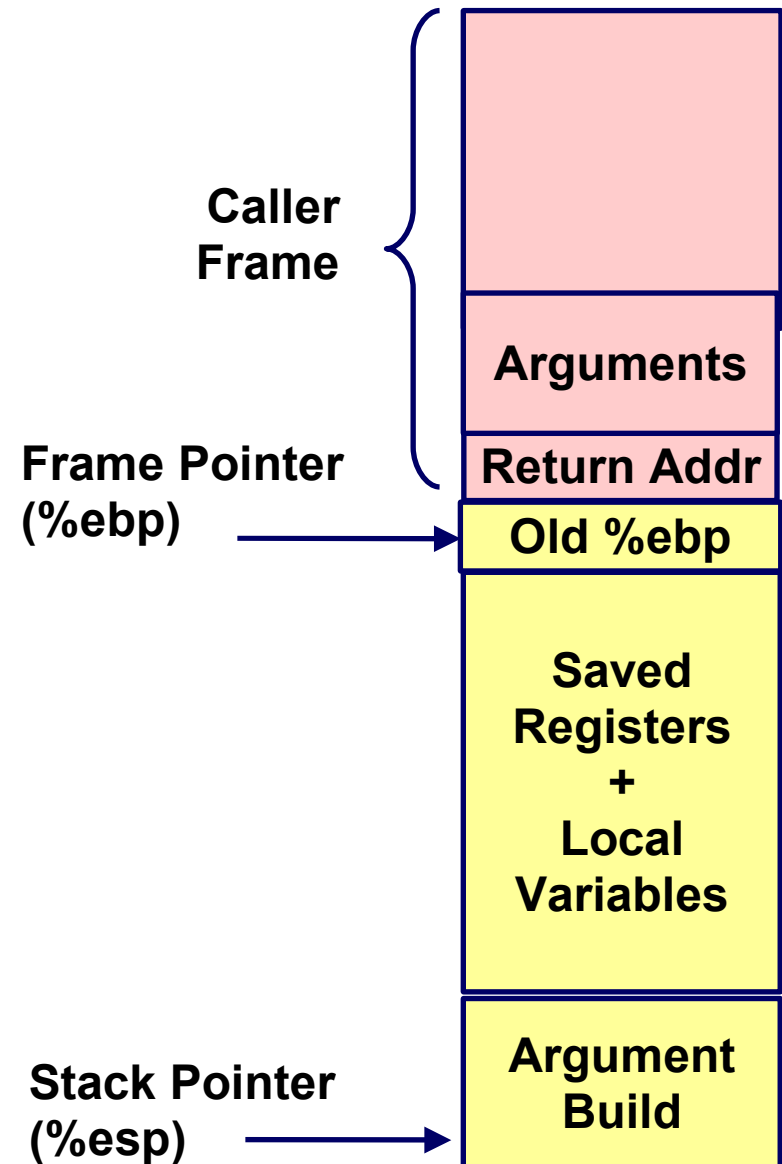
- Pointers

- Stack Pointer `%esp` 指向 Stack 的頂端
- Frame pointer `%ebp` 指向 Current Frame 的開端



Stack Frame - IA32/Linux

- Current Stack Frame
 - Argument Build
 - Parameters for function about to call
 - Local variables
 - Saved register context
 - Old frame pointer
- Caller Stack Frame
 - Return address
 - Pushed by **call** instruction
 - Arguments for this call

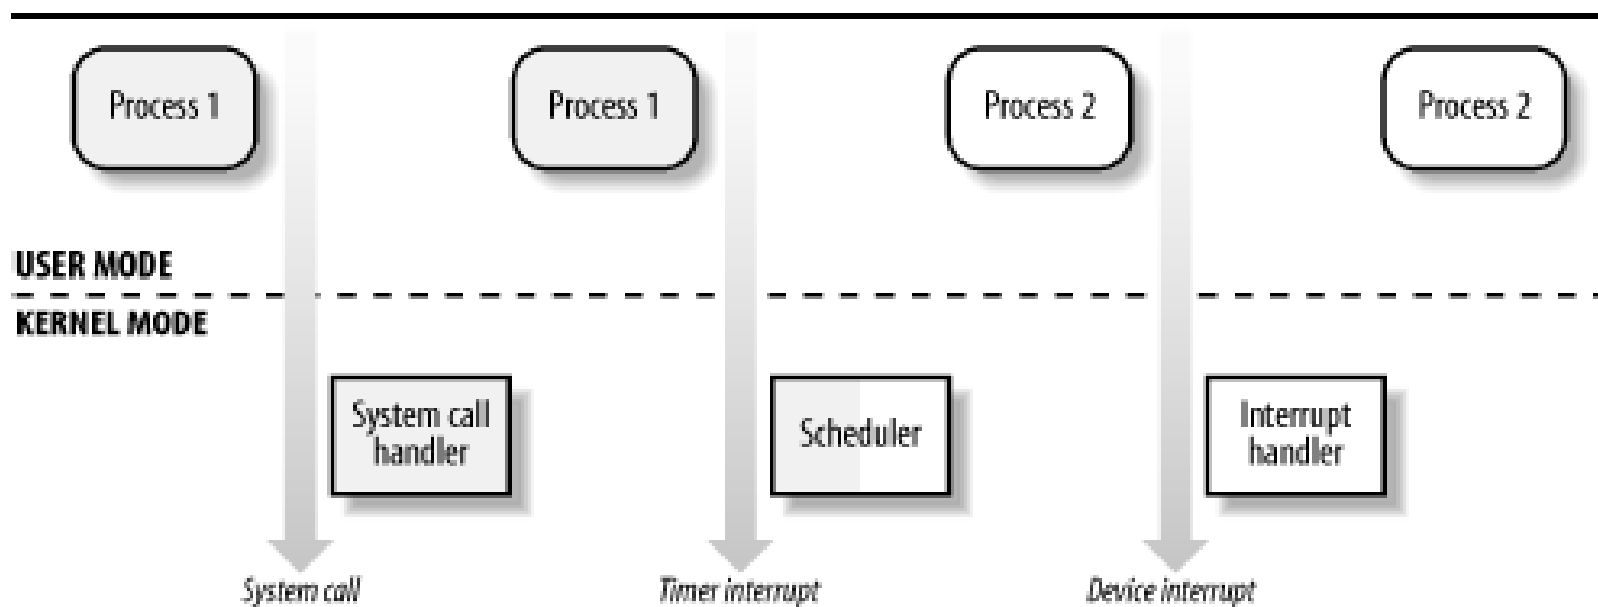
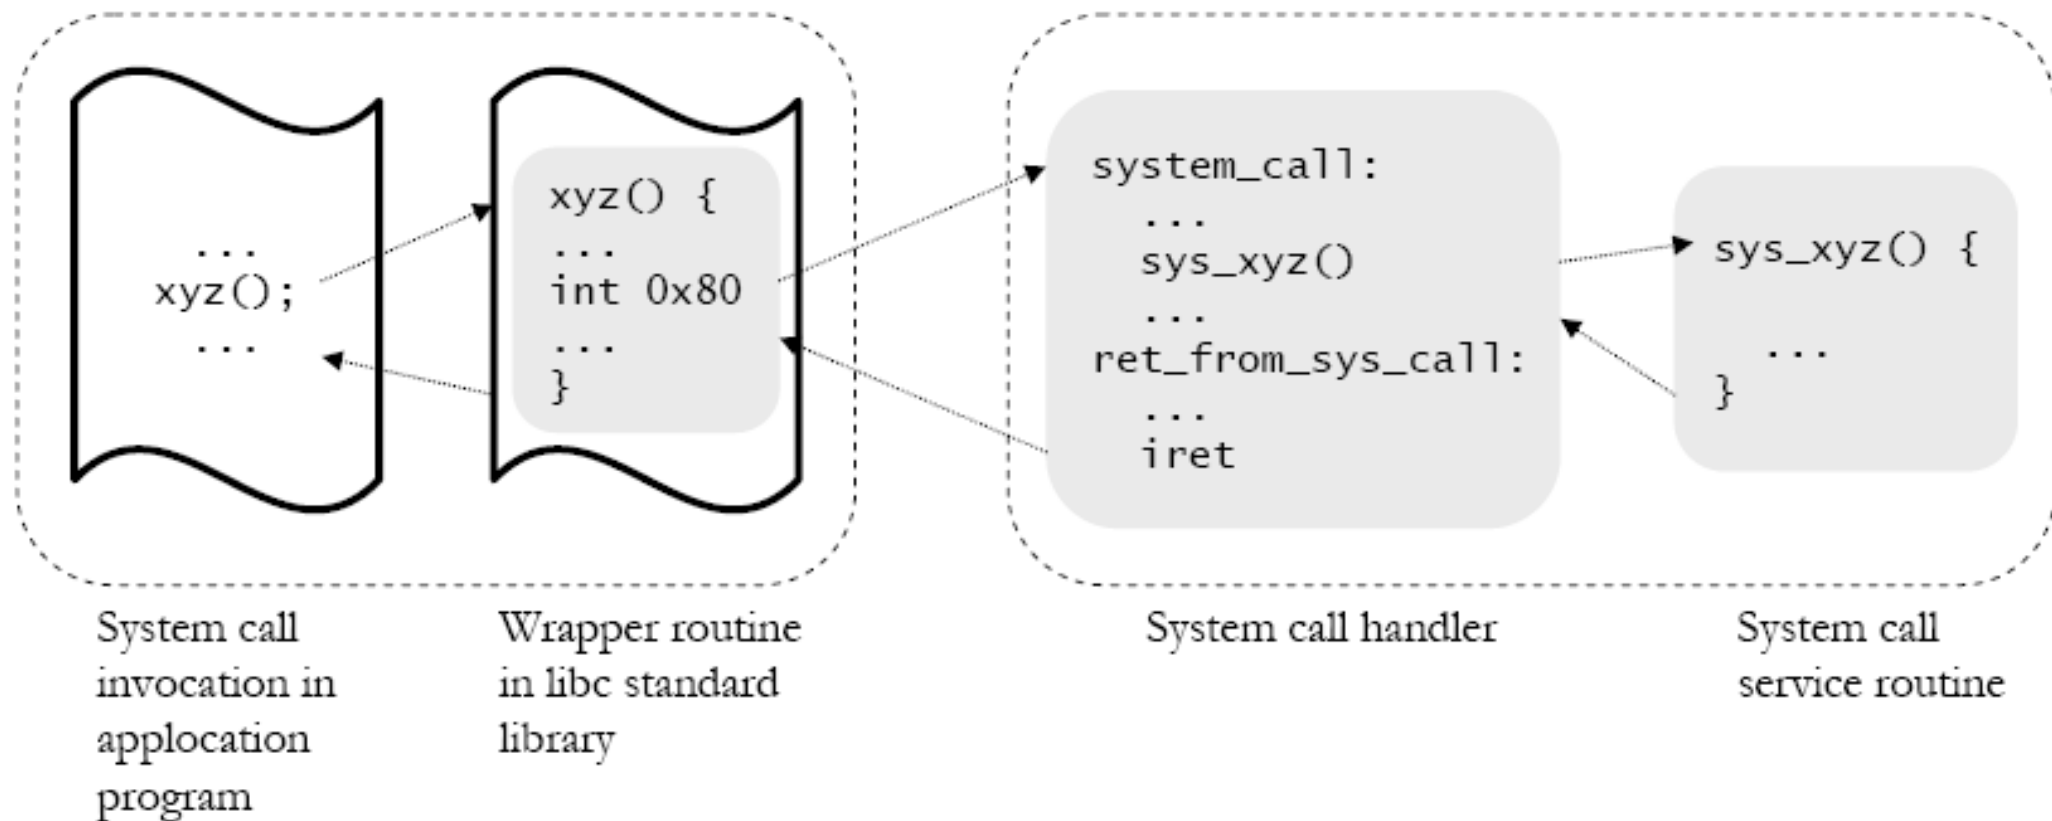


深入 syscall

- syscall = system call
- 機制
 - libc wrapper
 - direct syscall
 - syscall function
- Linux kernel 2.6 的 virtual syscall


User mode

Kernel mode




```
$ cat hello-loop.c
#include <stdio.h>
#include <unistd.h>
int main(int argc, char **argv)
{
    printf("Hello World!\n");
    while (1) {
        usleep(10000);
    }
    return 0;
}
$ ./hello-loop
Hello World!
```

```
$ pidof hello-loop
6987
$ gdb
(gdb) attach 6987
Attaching to process 6987
Reading symbols from
/home/jserv/HelloWorld/samples/hello-loop...done.
Using host libthread_db library
"/lib/tls/i686/cmov/libthread_db.so.1".
Reading symbols from
/lib/tls/i686/cmov/libc.so.6...done.
Loaded symbols for /lib/tls/i686/cmov/libc.so.6
Reading symbols from /lib/ld-linux.so.2...done.
Loaded symbols for /lib/ld-linux.so.2
0xfffffe410 in __kernel_vsyscall ()
```



```
(gdb) bt
#0  0xfffffe410 in __kernel_vsyscall ()
#1  0xb7e37ef0 in nanosleep () from /lib/tls/i686/cmov/libc.so.6
#2  0xb7e6f93a in usleep () from /lib/tls/i686/cmov/libc.so.6
#3  0x080483ad in main () at hello-loop.c:7
```

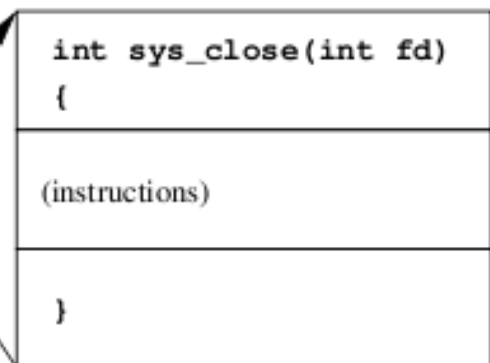

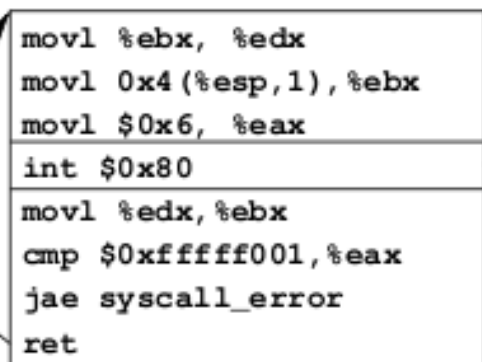

典型系統呼叫 (1)

main.c

(/lib/libc.so.6)

linux/arch/i386/kernel/entry.S

linux/fs/open.c



user mode | kernel mode

processes
not involved

system
calls

syscalls
not involved

assembler-glue
syscall multiplexer

user mode

kernel mode

典型系統呼叫 (2)

- 在 x86 保護模式，處理 `int` 中斷指令時，CPU 需要作繁複的查表與確認動作，才得以切換執行權，當 Kernel 執行系統呼叫完畢後，以 `iret` 返回，該指令恢復目前的 `stack`，並回到原本的執行
 - $CPL \leq DPL$
- 由 Ring3 進入 Ring0 的過程浪費許多 CPU 週期

回頭看 hello.c 的編譯過程

```
$ gcc -v -o hello hello.c
Using built-in specs.
Target: i486-linux-gnu
```

```
...
/usr/lib/gcc/i486-linux-gnu/4.1.2/collect2 --eh-frame-hdr
-m elf_i386 -dynamic-linker /lib/ld-linux.so.2 -o hello
/usr/lib/gcc/i486-linux-gnu/4.1.2/../../../../lib/crt1.o
/usr/lib/gcc/i486-linux-gnu/4.1.2/../../../../lib/crti.o
/usr/lib/gcc/i486-linux-gnu/4.1.2/crtbegin.o -L
/usr/lib/gcc/i486-linux-gnu/4.1.2 -L/usr/lib/gcc/i486-
linux-gnu/4.1.2 -L/usr/lib/gcc/i486-linux-
gnu/4.1.2/../../../../lib -L/lib/../../lib -L/usr/lib/../../lib
/tmp/ccyj1YoV.o -lgcc --as-needed -lgcc_s --no-as-
needed -lc -lgcc --as-needed -lgcc_s --no-as-needed
/usr/lib/gcc/i486-linux-gnu/4.1.2/crtend.o
/usr/lib/gcc/i486-linux-gnu/4.1.2/../../../../lib/crtn.o
$ wc -c hello
6749 hello
```

crt*.o

C Runtime object files

```
$ cat hello.c
```

```
int main() {
```

```
    printf("Hello World\n");
```

```
    return 0;
```

```
}
```

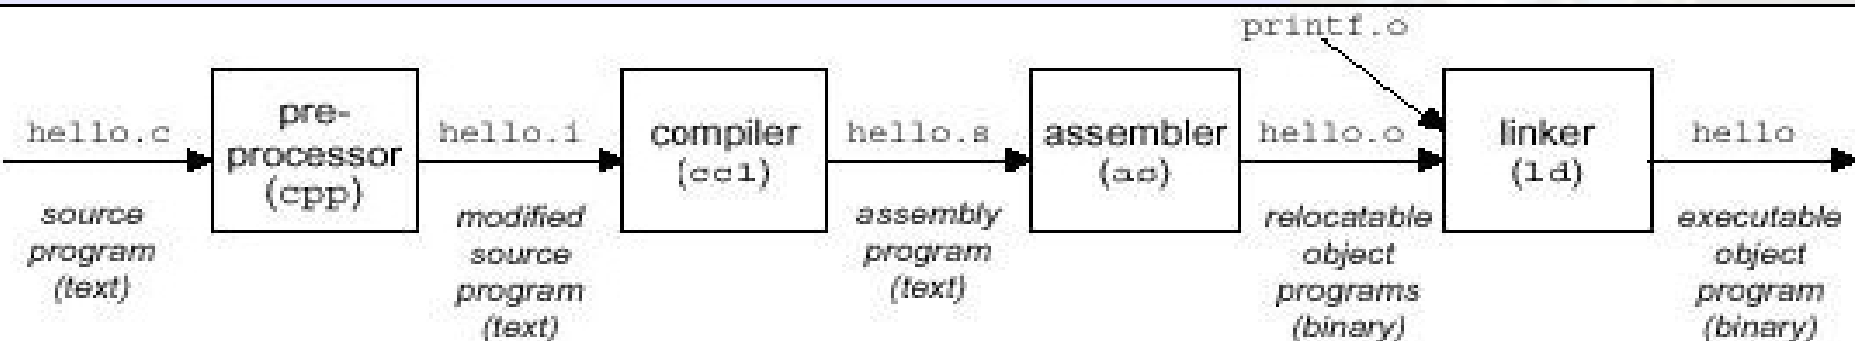
```
$ strace ./hello
```

```
...
```

```
write(1, "Hello World\n", 12Hello World.) = 12
```

```
exit_group(0)                                = ?
```

```
Process 14211 detached
```



系統呼叫 (1)

```
$ cat hello-write.c
#include <unistd.h>
int main()
{
    write(1, "Hello World\n", 12);
    return 0;
}
$ ./hello-write
Hello World
```

- 改用 POSIX 之 **write** 系統呼叫
- 但還不是最「直接」的使用系統呼叫
- 透過 **libc wrapper** 實現

/* Write N bytes of BUF to FD. Return the number written, or -1.

This function is a cancellation point and therefore not marked with
__THROW. */
extern ssize_t **write** (int __fd, __const void *__buf, size_t __n) __wur;

(/usr/include/unistd.h)

系統呼叫 (2)

```
$ cat hello-syscall.c
#include <asm/unistd.h>

static int errno;
__syscall1(int, exit, int, status);
__syscall3(int, write,
           int, fd,
           const void*, buf,
           unsigned long, count);

int main()
{
    write(1, "Hello World\n", 12);
    exit(0);
}
$ ./hello-syscall
Hello World
```

- linux-2.6.17/include/asm-i386/unistd.h
- gcc -o hello-syscall \
- D__KERNEL__** \
- fno-builtin \
- fomit-frame-pointer \
- I/lib/modules/`uname -r`/build/include \
- hello-syscall.c

```
#define __syscall1(type,name,type1,arg1) \
type name(type1 arg1) \
{ \
    long __res; \
    __asm__ volatile ("push %%ebx ; movl %2,%%ebx ; int \
                      $0x80 ; pop %%ebx" \
                      : "=a" (__res) \
                      : "0" (__NR_##name),"ri" ((long)(arg1)) : "memory"); \
    __syscall_return(type,__res); \
}
```

```
extern ssize_t write (int __fd, __const void *__buf, size_t __n) __wur;
```

(/usr/include/unistd.h)

系統呼叫 (3)

```
$ cat hello-syscall2.c
#include <asm/unistd.h>

static int errno;
_syscall1(int, exit, int, status);
_syscall3(int, write,
  int, fd, const void*, buf,
  unsigned long, count);

void hello()
{
  write(1, "Hello World\n", 12);
  exit(0);
}
```

```
$ gcc -c -Os -D__KERNEL__ \
  -fno-builtin -fomit-frame-pointer \
  -I/lib/modules/`uname -r`/build/include \
  hello-syscall2.c
$ ld -o hello-syscall2 \
  -e hello \
  hello-syscall2.o
$ wc -c hello-syscall2
977 hello-syscall2
```

```
#include <stdio.h>
char message[] = "Hello, world!\n";
int main(void) {
  long _res;
  __asm__ volatile (
    "int $0x80"
    : "=a" (__res)
    : "a" ((long) 4),
    "b" ((long) 1),
    "c" ((long) message),
    "d" ((long) sizeof(message)));
  return 0;
}
```

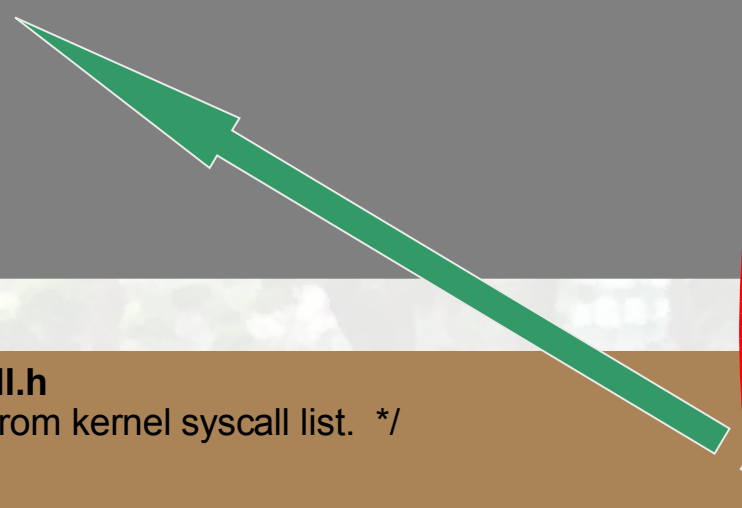
```
$ gcc -c -E -Os -D__KERNEL__ -fno-builtin -fomit-frame-pointer \
-I/lib/modules/`uname -r`/build/include hello-syscall2.c
```

```
...
int write(int fd,const void* buf,unsigned long count) { long __res; __asm__ volatile
("push %%ebx ; movl %2,%%ebx ; int $0x80 ; pop %%ebx" : "=a" (__res) : "0"
(4),"ri" ((long)(fd)),"c" ((long)(buf)), "d" ((long)(count)) : "memory"); do { if
((unsigned long)(__res) >= (unsigned long)(-(128 + 1))) { errno = -(__res); __res =
-1; } return (int) (__res); } while (0); };
...
```

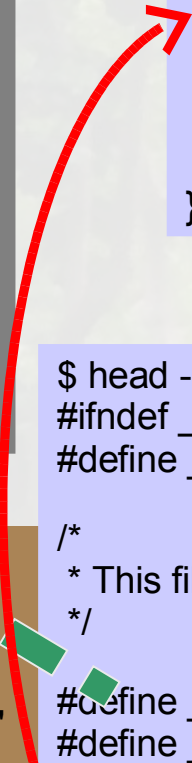

系統呼叫 (4)

```
$ cat hello-syscall3.c
#include <stdio.h>
#include <sys/syscall.h>
#include <unistd.h>

int main()
{
    int ret;
    ret = syscall(__NR_write, 1, "Hello World\n", 12);
    return 0;
}
$ ./hello-syscall3
Hello World
```



```
#include <stdio.h>
char message[] = "Hello, world!\n";
int main(void) {
    long _res;
    __asm__ volatile (
        "int $0x80"
        : "=a" (_res)
        : "a" ((long) 4),
          "b" ((long) 1),
          "c" ((long) message),
          "d" ((long) sizeof(message)));
    return 0;
}
```



```
$ head -n 12 /usr/include/asm-i386/unistd.h
#ifndef _ASM_I386_UNISTD_H_
#define _ASM_I386_UNISTD_H_

/*
 * This file contains the system call numbers.
 */

#define __NR_restart_syscall    0
#define __NR_exit                1
#define __NR_fork                2
#define __NR_read                3
#define __NR_write                4
```

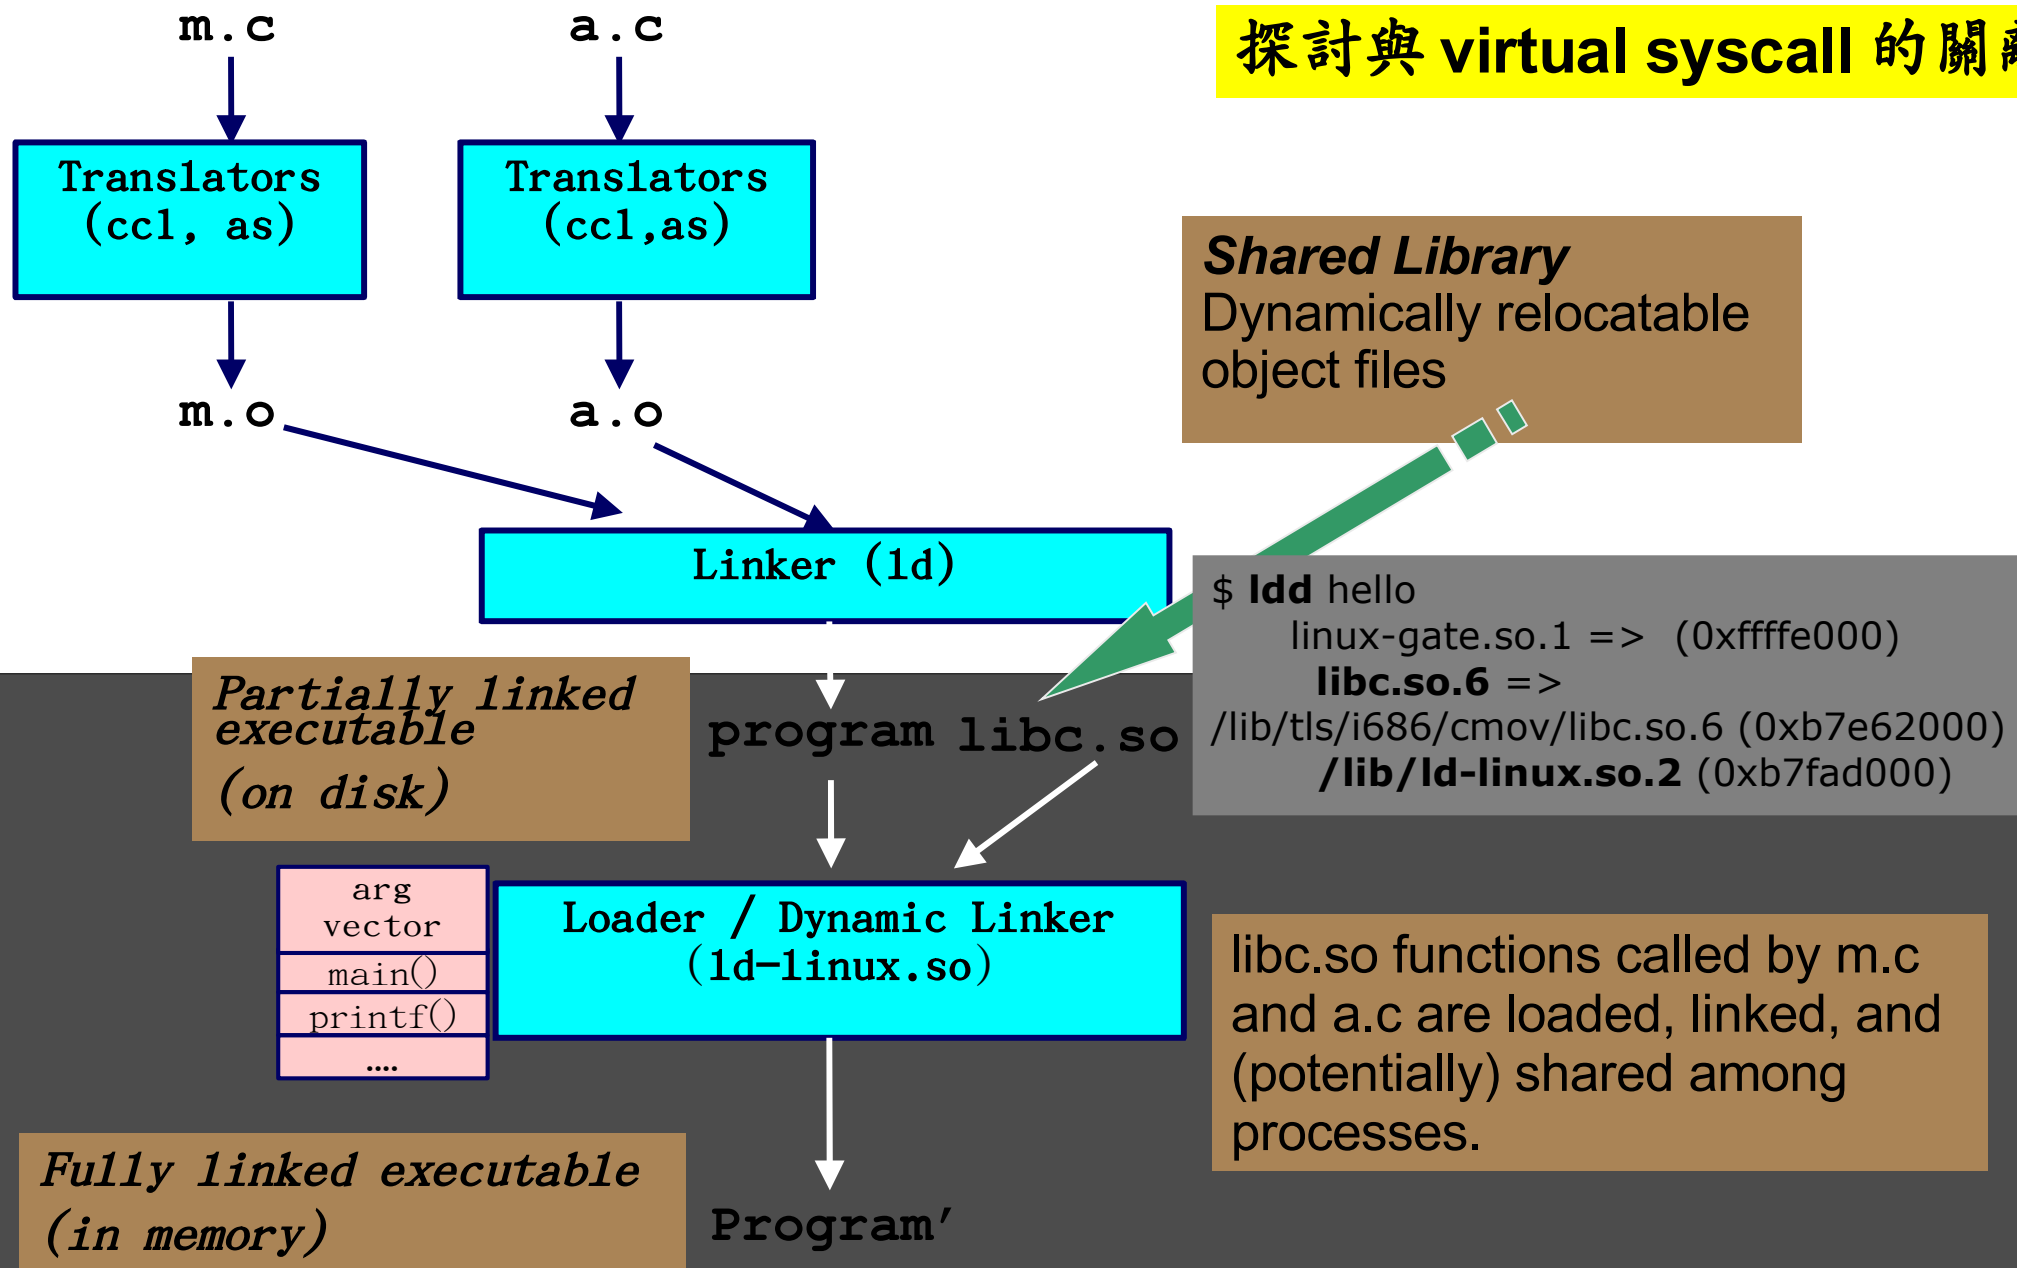
```
$ head /usr/include/bits/syscall.h
/* Generated at libc build time from kernel syscall list. */

#ifndef _SYSCALL_H
# error "Never use <bits/syscall.h> directly; include <sys/syscall.h> instead."
#endif

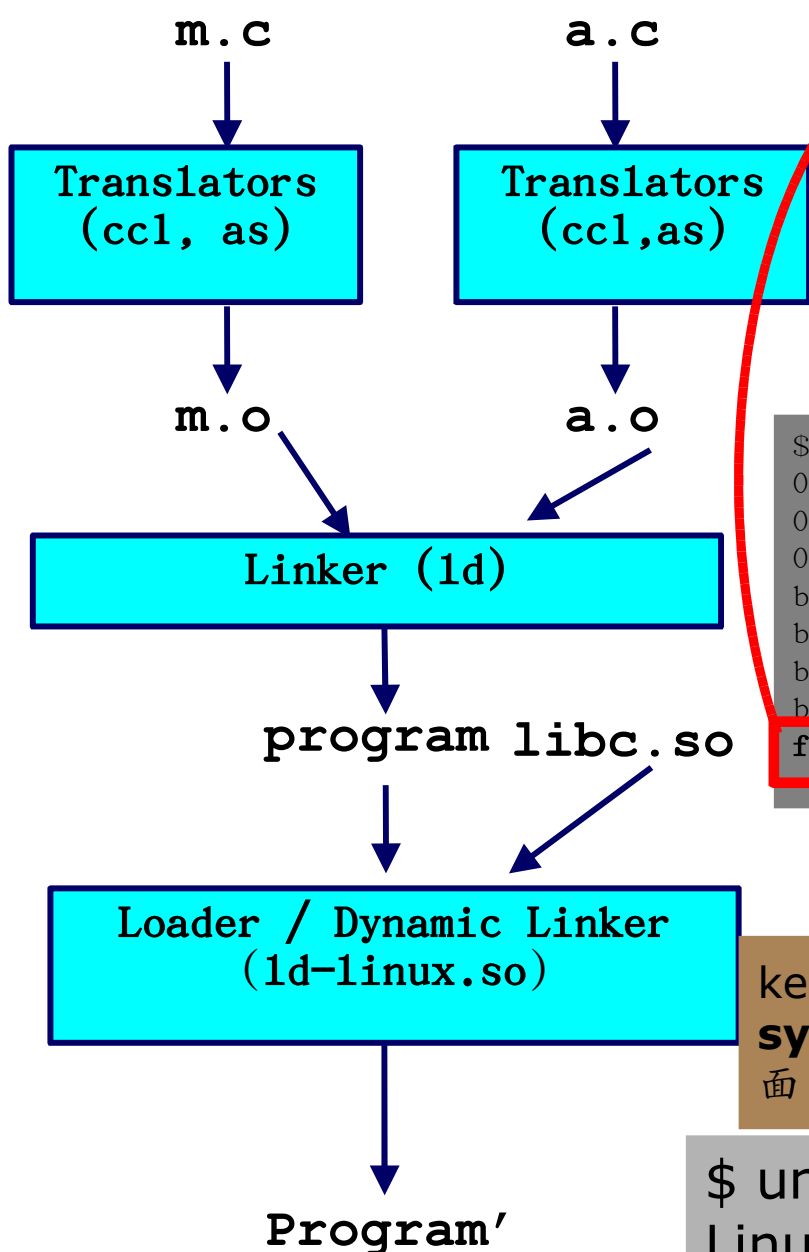
#define SYS__llseek __NR__llseek
#define SYS__newselect __NR__newselect
#define SYS__sysctl __NR__sysctl
#define SYS__access __NR__access
```


Dynamically Linked Shared Libraries⁽¹⁾

探討與 **virtual syscall** 的關離



Dynamically Linked Shared Libraries(2)



```
$ ldd hello
```

```
linux-gate.so.1 => (0xffffe000)
```

```
libc.so.6 =>
```

```
/lib/tls/i686/cmov/libc.so.6 (0xb7e62000)
```

```
/lib/ld-linux.so.2 (0xb7fad000)
```

/proc/self/maps
Memory map

```
$ cat /proc/self/maps
```

```
08048000-0804c000 r-xp 00000000 31222 /bin/cat
0804c000-0804d000 rw-p 00003000 31222 /bin/cat
0804d000-0806e000 rw-p 0804d000 0 [heap]
b7f90000-b7f92000 rw-p b7f90000 0
b7f92000-b7fab000 r-xp 00000000 295771 /lib/ld-2.4.so
b7fab000-b7fad000 rw-p 00018000 295771 /lib/ld-2.4.so
bfd95000-bfdab000 rw-p bfd95000 0 [stack]
ffffe000-fffff000 ---p 00000000 0 [vdso]
```

vdso : Virtual DSO

kernel 2.6 支援 Pentium III+ 的快速系統呼叫 **sysenter** / **sysexit** 機制。核心提供原先的 `int 0x80`，與新的 **sysenter** 雙介面，並虛擬印射了 **linux-gate.so.1** 來實作

```
$ uname -a
```

```
Linux venux 2.6.15-23-686 #1 SMP PREEMPT Tue
May 23 14:03:07 UTC 2006 i686 GNU/Linux
```

int80 vs. sysenter/sysexit

耗費時間

以 **int80** 為基礎的
系統呼叫

User-Mode

17.500ms

Kernel-Mode

7.00ms


Intel® Pentium® III CPU, 450 MHz Processor Family: 6
Model: 7 Stepping: 2

以 **sysenter/**
sysexit 為基礎

9.833ms

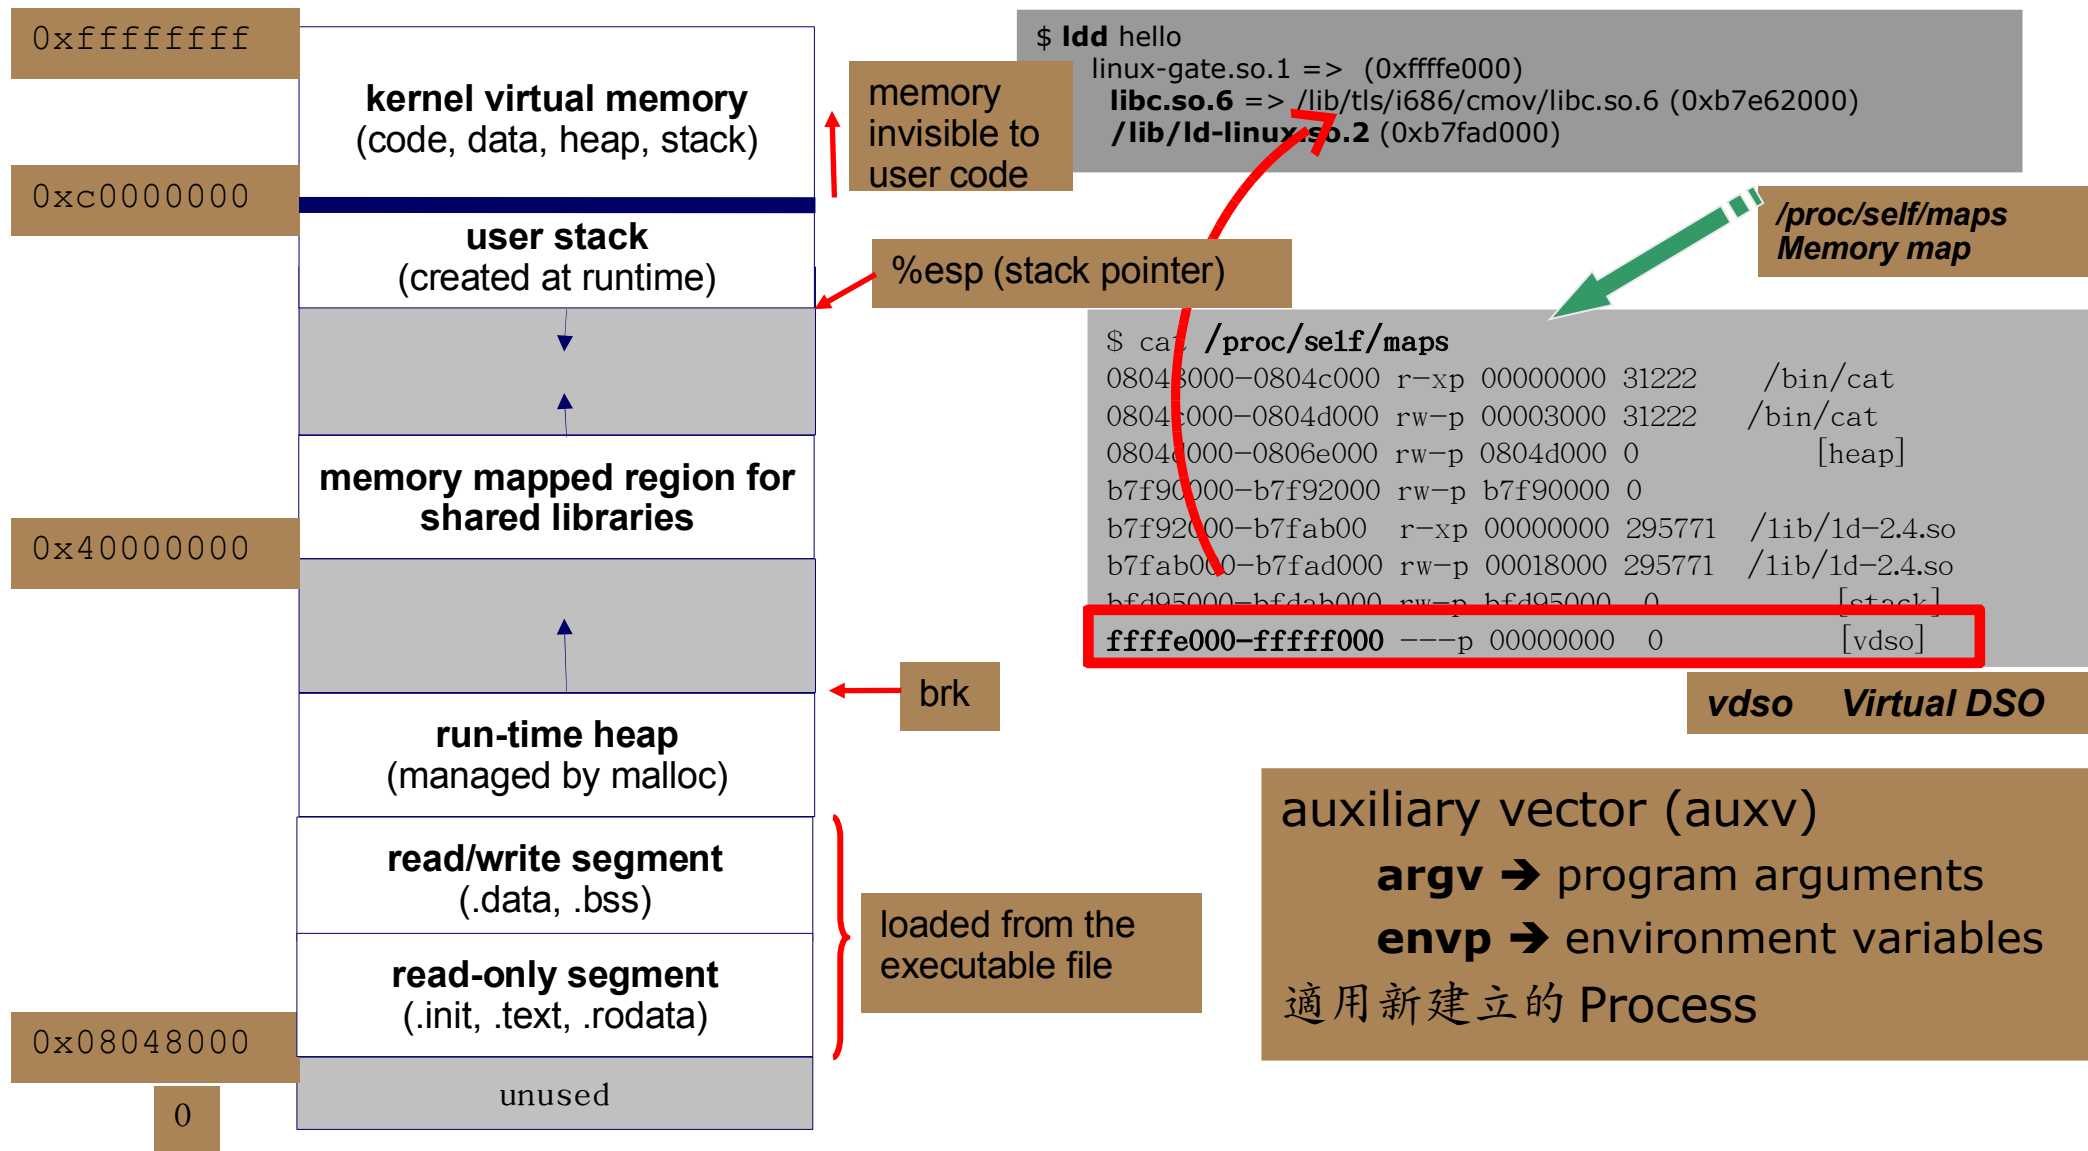
6.833ms

Performance System entry / exit from user level



on P III 500 Mhz

Dynamically Linked Shared Libraries(3)



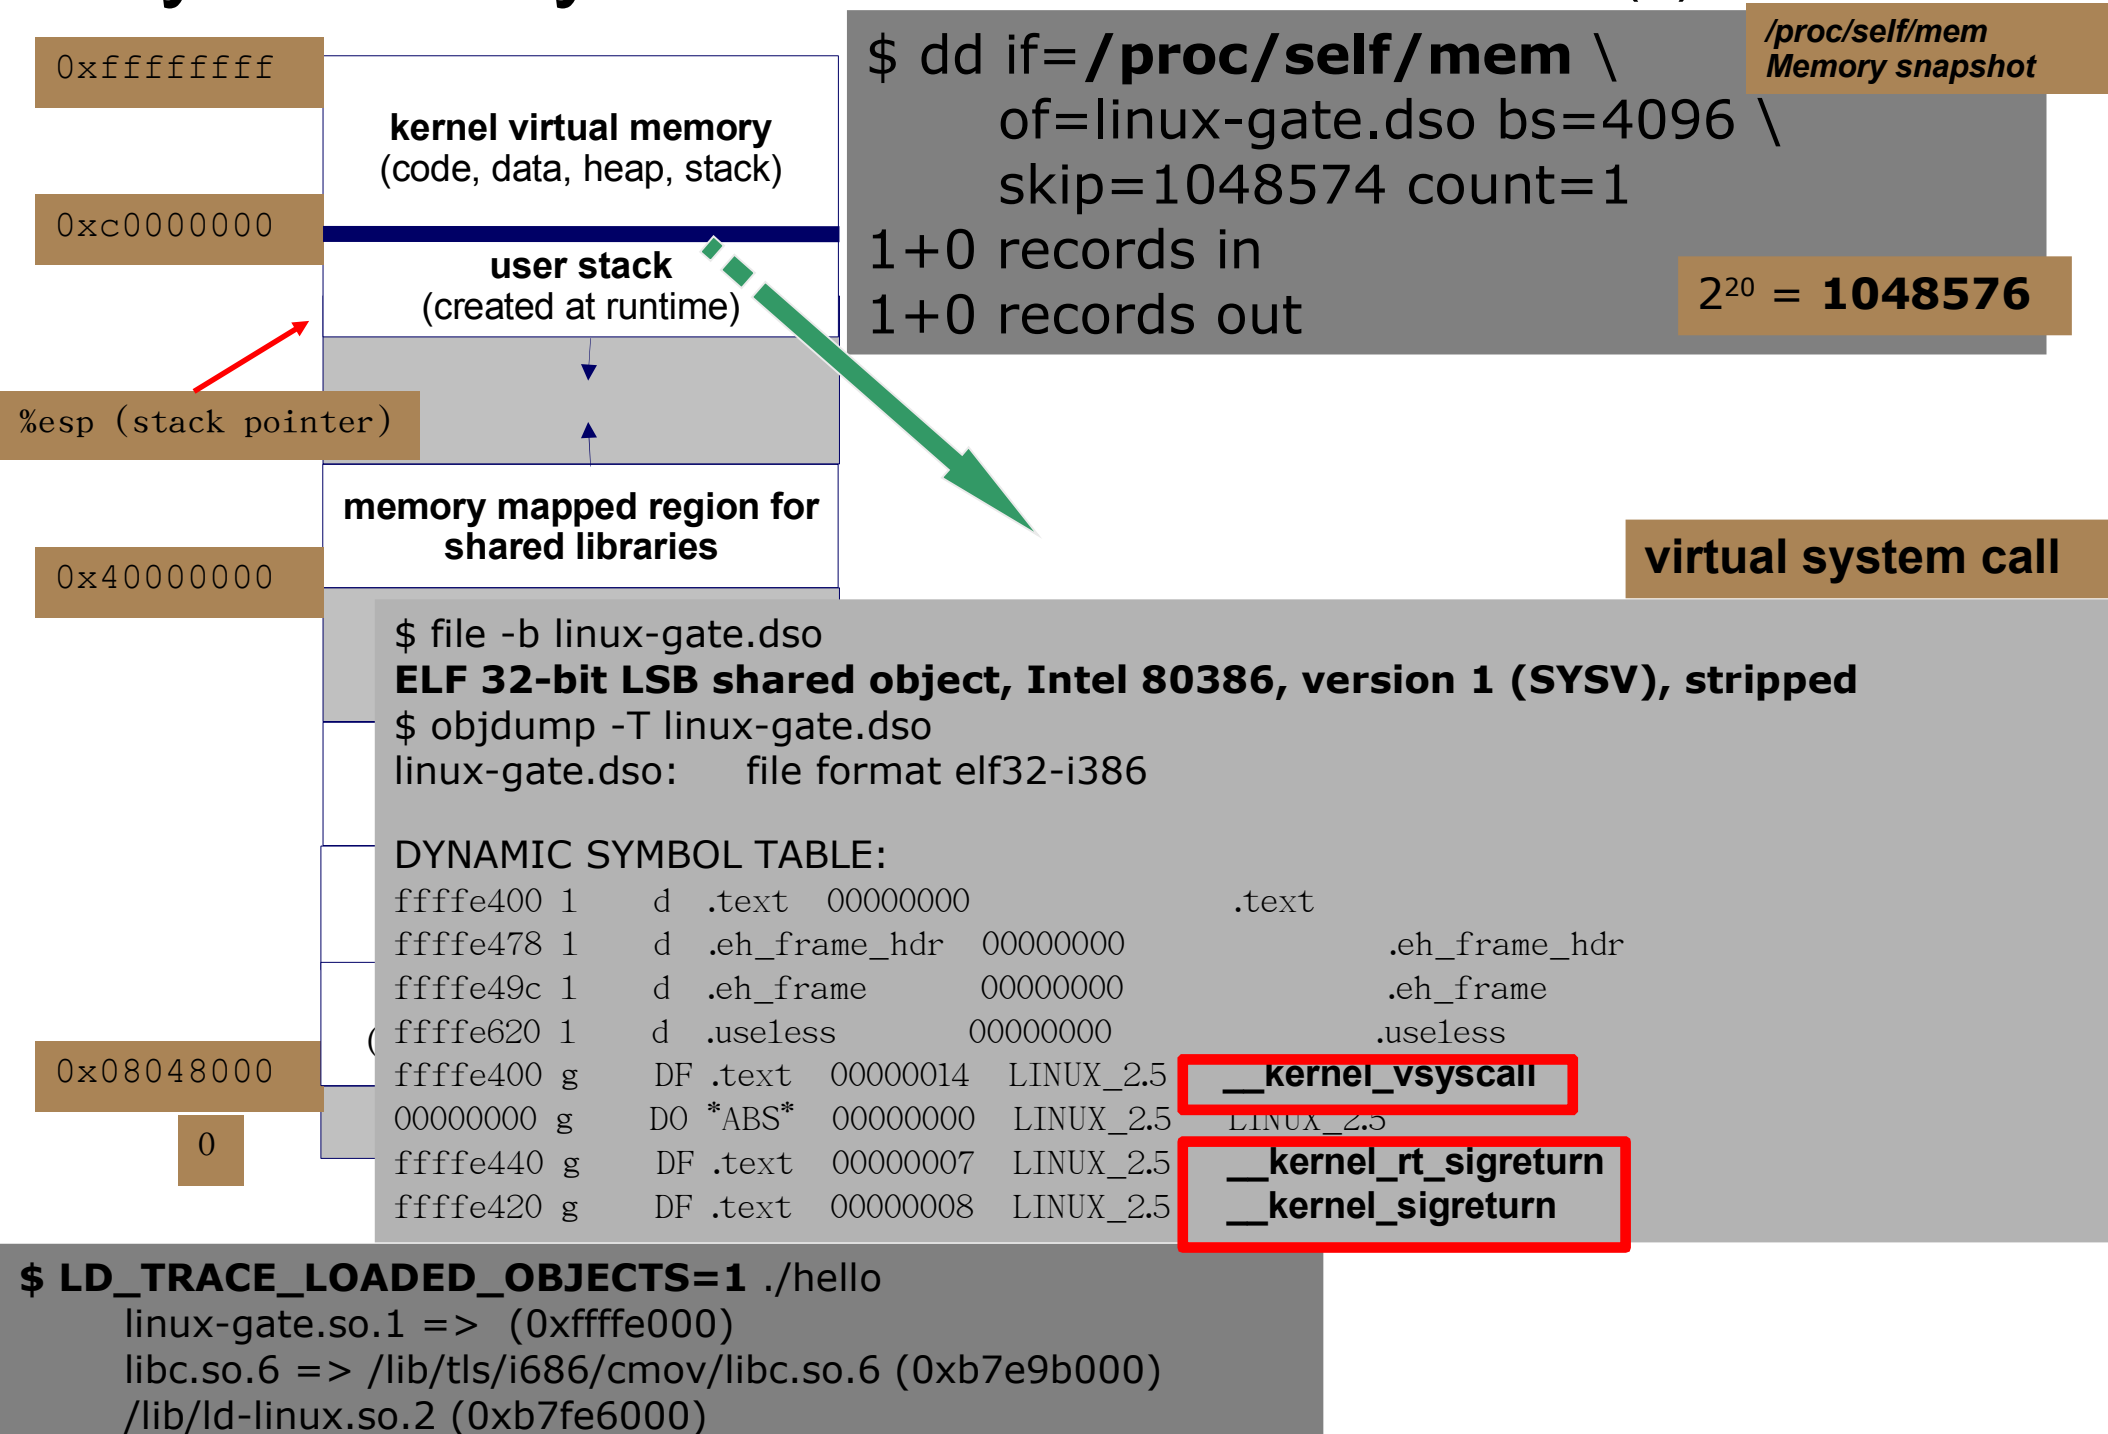
```
$ gdb -c core -q
Core was generated by `hello'.
Program terminated with signal 11, Segmentation fault.
#0 0xffffe777 in ?? ()
(gdb) bt
#0 0xffffe777 in ?? ()
```

VA stack exploit

Ref: Exploiting with linux-gate.so.1 . by Izik

```
# JMP *%ESP @ linux-gate.so.1jmp = "\x77\xe7\xff\xff"
```


Dynamically Linked Shared Libraries(4)



Dynamically Linked Shared Libraries(5)

0xffffffff

kernel virtual memory
(code, data, heap, stack)

virtual system call

0xc0000000

user
(created a)

```
$ file -b linux-gate.dso
```

ELF 32-bit LSB shared object, Intel 80386, version 1 (SYSV), stripped

```
$ objdump -T linux-gate.dso
```

linux-gate.dso: file format elf32-i386

%esp (stack pointer)

memory mapped
shared 1

DYNAMIC SYMBOL TABLE:

ffffe400	1	d	.text	00000000	.text
ffffe478	1	d	.eh_frame_hdr	00000000	.eh_frame_hdr
ffffe49c	1	d	.eh_frame	00000000	.eh_frame
ffffe620	1	d	.useless	00000000	.useless
ffffe400	g	DF	.text	00000014	LINUX_2.5
00000000	g	DO	*ABS*	00000000	LINUX_2.5
ffffe440	g	DF	.text	00000007	LINUX_2.5
ffffe420	g	DF	.text	00000008	LINUX_2.5

__kernel_vsyscall

__kernel_rt_sigreturn

__kernel_sigreturn

run-time
(managed by)

read/write segment
(.data, .bss)

--start-address=0xffffe400

__kernel_vsyscall 進入點

```
$ objdump -d --start-address=0xffffe400  
--stop-address=0xffffe414 linux-gate.dso
```

linux-gate.dso: file format elf32-i386
Disassembly of section .text:

ffffe400 <__kernel_vsyscall>:

ffffe400:	51	push	%ecx
ffffe401:	52	push	%edx
ffffe402:	55	push	%ebp
ffffe403:	89 e5	mov	%esp,%ebp
ffffe405:	0f 34	sysenter	

ffffe407:	50	nop	
ffffe408:	90	nop	
ffffe409:	90	nop	
ffffe40a:	90	nop	
ffffe40b:	90	nop	
ffffe40c:	90	nop	
ffffe40d:	90	nop	
ffffe40e:	eb f3	jmp	ffffe403
<__kernel_vsyscall+0x3>			
ffffe410:	5d	pop	%ebp
ffffe411:	5a	pop	%edx
ffffe412:	59	pop	%ecx
ffffe413:	c3	ret	

虛擬系統呼叫

```
$ cat hello.c
#include <stdio.h>
#include <sys/syscall.h>
#include <unistd.h>
int main() {
    int ret;
    char message[] = "Hello, world!\n";
    __asm__ volatile (
        "call *%2"
        : "=a" (ret)
        : "a" (__NR_write),
          "S" (0xffffe400),
          "b" ((long) 1),
          "c" ((long) message),
          "d" ((long) sizeof(message)));
    return 0;
}
$ ./hello
Hello, world!
```

```
$ objdump -d --start-address=0xffffe400 \
--stop-address=0xffffe414 linux-gate.dso

linux-gate.dso:      file format elf32-i386
Disassembly of section .text:

fffffe400 <_kernel_vsyscall>:
fffffe400:    51          push    %ecx
fffffe401:    52          push    %edx
fffffe402:    55          push    %ebp
fffffe403:    89 e5       mov     %esp,%ebp
fffffe405:    0f 34       sysenter
```

以 **procedure call** 的形式
呼叫了 **system call**

再探動態連結

- 執行時期的系統呼叫
- glibc 的 HWCAP 機制
- Rundll32.exe on Linux


```

[x]— /home/jserv/HelloWorld/helloworld/samples/00-pureC/hello
* ELF section headers at offset 000007e4
[+] section 0:
[-] section 1: .interp
    name string index          0000000b
    type                      00000001 (progbits)
    flags                     00000002 details
    address                   08048114
    offset                    00000114
    size                      00000013
    link                      00000000
    info                      00000000
    alignment                 00000001
    entsize                   00000000
[+] section 2: .note.ABI-tag
[+] section 3: .hash
[+] section 4: .dynsym
[+] section 5: .dynstr
[+] section 6: .gnu.version
[+] section 7: .gnu.version_r
[+] section 8: .rel.dyn
[+] section 9: .rel.plt

```

The diagram illustrates the process of loading a shared library. A green arrow points from the `.interp` section header in the ELF file to the `elf_interpreter` field in the `ld-linux.so.2` file. Another green arrow points from the `elf_interpreter` field to the `/lib/ld-linux.so.2` file, indicating that the program loader will use this file to load the shared libraries needed by the program.

`.interp` →
`elf_interpreter`

\$ /lib/ld-linux.so.2

Usage: ld.so [OPTION]... EXECUTABLE-FILE [ARGS-FOR-PROGRAM...]

You have invoked `ld.so', the helper program for shared library executables. This program usually lives in the file `/lib/ld.so', and special directives in executable files using ELF shared libraries tell the system's program loader to load the helper program from this file. This helper program loads the shared libraries needed by the program executable, prepares the program to run, and runs it.

```
$ objdump -s -j .interp hello
```

```
hello:      file format elf32-i386
```

```
Contents of section .interp:
```

```

8048114 2f6c6962 2f6c642d 6c696e75 782e736f /lib/ld-linux.so
8048124 2e3200      2.
```

執行時期的系統呼叫 (1)

- glibc 的實做 sysdeps/unix/sysv/linux/i386/sysdep.h

```
# define INTERNAL_SYSCALL(name, err, nr, args...) \
```

```
{ \
    register unsigned int resultvar; \
    EXTRAVAR_##nr \
    asm volatile ( \
        LOADARGS_##nr \
        "movl %1, %%eax\n\t" \
        "int $0x80\n\t" \
        RESTOREARGS_##nr \
        : "=a" (resultvar) \
        : "i" (__NR_##name) \
        ASMFMT_##nr(args) \
        : "memory", "cc"); \
    (int) resultvar; }
```

```
#define INLINE_SYSCALL(name, nr, args...) \
({ \
    unsigned int resultvar = INTERNAL_SYSCALL \
    (name, , nr, args); \
    if (__builtin_expect \
    (INTERNAL_SYSCALL_ERROR_P (resultvar, ), 0)) \
    { \
        __set_errno (INTERNAL_SYSCALL_ERRNO \
    (resultvar, )); \
        resultvar = 0xffffffff; \
    } \
    } \
    (int) resultvar; })
```

執行時期的系統呼叫 (2)

- 程式載入器或 shell 會有類似的操作

- `execve` syscall

- Linux Kernel

- `arch/i386/kernel/traps.c`

```
void __init trap_init(void)
{
```

```
...
```

```
    set_system_gate(SYSCALL_VECTOR,
                    &system_call);
```

```
}
```

```
movl <envp>, %edx
```

```
movl <argv>, %ecx
```

```
movl <file>, %ebx
```

```
movl $11, %eax
```

```
int $0x80
```

```
; execve
```

將於 Part III 中透過 User-Mode-Linux 與 qemu 分析其原理

glibc 的 HWCAP 機制 (1)

- \$ COLUMNS=200 dpkg -l | grep libc6
 - ii **libc6** 2.5-0ubuntu2 GNU C Library:
Shared libraries
 - ii libc6-dev 2.5-0ubuntu2 GNU C
Library: Development Libraries and Header Files
 - ii **libc6-i686** 2.5-0ubuntu2 GNU C
Library: Shared libraries [**i686 optimized**]
- 一般 i386 與 i686 最佳化的 glibc 如何共存？
- 特定之函數，如數學運算，如何在執行時期針對硬體挑選最佳的實做？
- 引入 HWCAP (Hardware Capacities) 的機制
 - LD_SHOW_AUXV (AUXiliary Vector)

glibc 的 HWCAP 機制 (2)

```
$ ldd hello
linux-gate.so.1 => (0xffffe000)
libc.so.6 => /lib/tls/i686/cmov/libc.so.6 (0xb7e8d000)
/lib/ld-linux.so.2 (0xb7fe5000)
$ realpath /lib/libc.so.6
/lib/libc-2.5.so
$ strace -f ./hello
execve("./hello", ["./hello"], [/ * 28 vars */]) = 0
...
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or
directory)
open("/lib/tls/i686/cmov/libc.so.6", O_RDONLY) = 3
read(3, "\177ELF\1\1\1\0\0\0\0\0\0\0\0\0\3\0\3\0\1\0\0\0\0\1\000"..., 512) = 512
fstat64(3, {st_mode=S_IFREG|0644, st_size=1311200, ...}) = 0
...
```

cmov = Pentium4 SSE2 Conditional MOVe

glibc 的 HWCAP 機制 (3)

```
$ LD_SHOW_AUXV=1 ./hello
```

```
AT_SYSINFO: 0xffffe400
```

```
AT_SYSINFO_EHDR: 0xffffe000
```

```
AT_HWCAP: fpu vme de pse tsc msr mce cx8 sep mtrr pge mca cmov pat  
clflush dts acpi mmx fxsr sse sse2 tm pbe
```

```
AT_PAGESZ: 4096
```

```
AT_CLKTCK: 100
```

```
AT_PHDR: 0x8048034
```

```
AT_PHEENT: 32
```

```
AT_PHNUM: 7
```

```
AT_BASE: 0xb7fca000
```

```
AT_FLAGS: 0x0
```

```
AT_ENTRY: 0x80482b0
```

```
AT_UID: 1000
```

```
AT_EUID: 1000
```

```
AT_GID: 1000
```

```
AT_EGID: 1000
```

```
AT_SECURE: 0
```

```
AT_PLATFORM: i686
```

```
Hello World!
```

Rundll32.exe on Linux⁽¹⁾

- MS-Windows Rundll.exe 允許透過命令列載入 DLL 並呼叫其中的 function
 - Rundll32.exe DllFileName FuncName
- izik(<http://www.tty64.org>) 針對 x86/Linux 撰寫 Runlib32
 - ./runlib libc.so.6,puts \'"Hello World"\'

```
$ ./runlib -v -x printf-out libc.so.6,puts \'"Hello World"\'  
puts[<0xb7ed8610>]@libc.so.6[]
```

```
-----  
* Stack Generated (1 parameters, 4 bytes)  
-----
```

Generated Assembly

```
-----  
* pushl $0xbfce7c9a  
* call 0xb7ed8610
```

Streams Buffers

```
-----  
* Standart Output (STDOUT) : 15 bytes  
* Standart Error (STDERR) : 0 bytes
```

Function Result

```
-----  
* Pointer: No  
* Value: 12  
$ cat printf-out  
Hello World
```

Rundll32.exe on Linux(2)

src/lib.c


```
/*
 * Manually pushing the function arguments to
 * the stack
 */

if (ptr->stack) {
    for (j = 0; j < ptr->stack->stack_items; j++) {
        __asm__ __volatile__ (\
            "pushl %0 \n" \
            : /* no output */ \
            : "r" (ptr->stack->stack[j]) \
            : "%eax" \
        );
    }
}
```

```
/*
 * Make the CALL!
 */
ret = (unsigned long) ptr->fcn_handler();

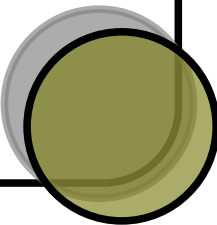
/*
 * Be polite, let's clean the stack afterward
 */
if (ptr->stack) {
    ptr->stack->stack_items *= sizeof(long);
    __asm__ __volatile__ (\
        "addl %0, %%esp \n" \
        : /* no output */ \
        : "r" (ptr->stack->stack_items) \
        : "%esp" );
    ptr->stack->stack_items /= sizeof(long);
}
```

```
s_errno = errno;
signal(SIGSEGV, SIG_DFL);
```



Rundll32.exe on Linux⁽³⁾

- Trampoline (Assembly/Machine_code-C interfacing)
- 類似技術廣泛應用於：
 - ffcall (GNUstep/Objective-C)
 - libffi (GNU GCC)
 - JIT compiler (Kaffe, Hotspot, ...)
 - Boot-strapping code in dynamic programming language engine



參考資料

- 《 Binary Hacks 》 , O'Reilly Japan
- Linkers and Loaders
 - <http://www.iecc.com/linker/>
- Startup state of a Linux/i386 ELF binary
 - <http://asm.sourceforge.net/articles/startup.html>
- IA32 上 Linux 內核中斷機制分析
 - <http://www.whitecell.org/list.php?id=23>
- FFCALL
 - <http://www.haible.de/bruno/packages-ffcall.html>
- Linux 2.6 對新型 CPU 快速系統調用的支持
 - <http://www-128.ibm.com/developerworks/cn/linux/kernel/l-k26ncpu/>

Incoming Part III

- 以 User-Mode-Linux 與 qemu 分析系統呼叫
- 探索 Linux Kernel 之 Program Loader
- User-space 與 Kernel-space 的互動
- 效能與記憶體窺探

